

□ YATRA.AI

AUTONOMOUS TRAVEL OPERATING SYSTEM

PRODUCT MANAGEMENT CASE STUDY

AI-Native Travel Planning

From fragmented group travel planning to an intelligent,
collaborative travel operating system.

Product Strategy · Multi-Agent AI · Prototype · Validation · GTM

Author: Vijetha Nissankara
Product Strategy & Discovery Lead
BITSOM Product Management Capstone

PRODUCT MANAGEMENT CASE STUDY

YatraAI — AI-Native Travel Planning

From fragmented group travel planning to an intelligent, collaborative travel operating system.

Field	Summary
Thesis	Group travel planning is fragmented across 7+ disconnected apps. YatraAI is a multi-agent AI layer that coordinates research, consensus, budgeting, booking and post-trip support inside a single WhatsApp-first flow.
Core problem	How can travel planning be simplified through a multi-agent system that handles research, bookings, budget optimization and contextual recommendations?
Primary user	Urban Group Traveller (22–35, Tier-1 metro) — entry persona: Aarav Sharma, “The Reluctant Trip Admin”.
Key differentiator	A Group Consensus Engine — the one capability neither MakeMyTrip's Myra nor ixigo's TARA offer; both are single-traveller AI assistants.
My role	Product Strategy & Discovery, PRD creation, Feature Breakdown, User Stories, GTM Strategy & Visuals, User Journey, Jira Board, end-to-end report structuring.

► The product journey, end to end



THE PROBLEM

Every stage of group travel planning happens in a different app — and the coordination tax falls on one person.



None of these tools talk to each other. The organizer becomes the integration layer — manually comparing prices across apps, chasing friends for decisions and money, and reconciling expenses after the trip ends.

► Evidence the fragmentation is real, not assumed

- **OTA trust is low:** review analysis of MakeMyTrip and Goibibo shows an average rating of roughly 1.4–1.5 stars — a structural trust gap, not a one-off complaint.
- **No group tooling exists:** secondary research found zero dedicated group-consensus tooling anywhere in the current market.
- **Five recurring pain points surfaced in research:** too many apps required, group decision fatigue, refund anxiety, budget conflicts, and organizer overload — one person bearing most of the planning effort.



YATRAAI OPPORTUNITY

Turn fragmented coordination into one intelligent journey — research, consensus, budgeting, booking and post-trip support inside a single AI-orchestrated, WhatsApp-first flow.

RESEARCH → INSIGHTS

Secondary research shaped the problem statement; four primary experiments validated it with real participants.

Research type	Sources / method	What it established
Secondary	OTA reviews (MakeMyTrip, Goibibo), Reddit, travel forums, complaint portals, competitor analysis	The trust gap and the absence of group-consensus tooling in market
Primary — Task 01	5 JTBD interviews (video-recorded, 1:1)	Validated renewal-relevant and adoption-relevant traveller pain points
Primary — Task 02	5-participant, within-subject, two-channel (WhatsApp vs PWA) onboarding experiment	Which entry channel removes friction, not which one demo better
Primary — Task 04	3 business-owner interviews (fashion D2C, vehicle rental, home catering), 22–31 min each	Cross-industry support patterns applicable to an AI travel concierge

► **Participant mix across the experimental tasks**

- Group traveller
- Family traveller
- Leisure / business (bleisure) traveller
- Solo traveller
- Corporate traveller

► **Strongest recurring findings**

Too many apps required		5/5
Group decision fatigue		4/5
Refund anxiety		4/5
Budget conflicts		4/5
Organizer overload		5/5

PERSONAS & TARGET SEGMENT

Every persona's frustration map to a named YatraAI agent or module — but they are not equal priorities.

**PRIMARY PERSONA — AARAV SHARMA, “THE RELUCTANT TRIP ADMIN”**

Bengaluru · Software Engineer · unofficial trip planner-in-chief. “I want to enjoy the trip too, not spend weeks planning it.” Pain: chasing friends for decisions and money. Needs: automatic consensus, transparent shared budgeting, reliable booking. YatraAI capability: Consensus Agent + Shared Trip Pool + Autonomous Booking.

► Secondary segment — same Urban Group Traveller profile**Priya Mohan — Budget-Conscious Traveler**

Hyderabad, marketing executive. Pain: hidden fees, unexpected costs. Needs upfront total cost. → Budget Dashboard, UPI Trip Pool, Expense Prediction.

**Rohan Gupta — The Busy Professional**

Pune, product manager. Pain: research overload, too many apps. Needs one-click booking. → Research Agent, Booking Agent, WhatsApp Interface.

► Tertiary segments — sequenced for Release 2.0**Sneha Menon — The Family Traveler**

Noida, HR manager, plans for six. Pain: scattered documents, elderly coordination. → Document Vault, Consensus Agent, Support Agent.

**Kabir Singh — The Frequent Traveler**

Bhopal, management consultant. Pain: re-entering preferences every trip. → Memory Layer, Context Engine, AI Learning.

This prioritization mirrors the STP targeting decision in Chapter 16 of the source document: Urban Group Travelers first, Tier-2/3 first-time flyers sequenced for Release 2.0, and Bleisure/Corporate segments deliberately held back since they require a different B2B go-to-market motion.

PRODUCT OPPORTUNITY & VALUE PROPOSITION

Current state	YatraAI future state
Fragmented research across OTAs and review sites	AI-led discovery (Research Agent)
Group arguments over destination and dates	Consensus Engine with anonymous voting
Multiple booking platforms, manual comparison	Coordinated booking flow (human-approved)
Budget spreadsheets, surprise costs	Budget intelligence — upfront cost estimate
Manual post-trip settlement (Splitwise, cash)	Shared UPI Trip Pool
Reactive support — you have to ask	Context-aware travel companion (proactive)

"Plan Together. Travel Smarter."

YatraAI positioning statement

Positioning: for urban Indian friend groups who travel together, YatraAI is the AI-native travel platform that turns group trip planning from a WhatsApp coordination nightmare into a single guided conversation — unlike Myra and TARA, which are AI-powered but built for one traveller at a time, YatraAI is built around the group as the unit of decision-making from the ground up.

PRODUCT EXPERIENCE — THE END-TO-END JOURNEY

The journey is a closed loop, not a funnel — stage 8 feeds directly back into stage 1.

- 01 Trip created**
Group trip started via WhatsApp link/QR — no new app to install. [User action]
- 02 Preferences captured**
AI asks each member a few conversational questions instead of manual polling. [AI action]
- 03 Consensus reached**
Consensus Agent surfaces agreement/disagreement and returns a consensus score. [AI action]
- 04 Itinerary generated**
Research Agent proposes destination, flights, hotels with visible reasoning. [AI action]
- 05 Budget checked**
Total cost per person shown upfront, before anyone commits. [AI action]
- 06 Booking approved**
Booking executes only after explicit human approval; funds held in escrow-style pool until then. [Human approval gate]
- 07 In-trip support**
Support & Moments agents handle delays, rebooking, proactive nudges. [AI action]
- 08 Trip memory saved**
Memory Agent logs the trip and prompts the next one — looping back to stage 1. [AI action]

PRODUCT ARCHITECTURE — THE FEATURE ECOSYSTEM

Every module resolves a named pain point from a named persona — grouped by the role it plays in the journey.

 Entry & Collaboration Yatra Connect · Yatra Sync · Yatra Consensus	 Discovery & Planning Yatra Discover · Yatra Planner
 Transaction & Money Yatra Book · Yatra Budget · Yatra Wallet	 Travel Experience Yatra Vault · Yatra Rescue · Yatra Moments · Yatra Companion
 Memory & Intelligence Yatra Memories · Yatra Insights	 Operations Yatra Command (internal admin — user, booking, payment & fraud monitoring)



WORTH DEMOING FIRST

Yatra Moments, Yatra Companion and Yatra Memories are what make YatraAI feel like a companion rather than a booking form — the modules 12–14 in the source document.

MVP & PRIORITISATION



RELEASE 1.0 HYPOTHESIS

Can AI-mediated consensus and transparent, upfront budgeting get a group from intent to a booked trip — without full autonomous booking or live payments yet? Release 1.0 is deliberately scoped to test that, not to ship every capability.

Must Have (Release 1.0)	Why
User login & authentication	Secure onboarding, personalized experience
Trip creation & group invitations	Core entry point
AI preference collection	Conversational, not form-based
Group Consensus Engine (Yatra Sync)	Resolves conflicts into one accepted plan
AI Research Agent (Yatra Discover)	Replaces manual destination research
AI Itinerary Generator	Optimized day-wise plan from group preferences
Budget Dashboard (Yatra Wallet)	Real-time visibility and forecasting
UPI Trip Pool (prototype)	Demonstrates shared-wallet workflow
Basic booking flow (mock APIs)	Simulates the booking journey for MVP validation
Clickable prototype & wireframes	Usability testing before development

► Deliberately not built yet

Release 2.0 (should/could-have): Autonomous Booking Agent, Live Price Monitoring, Travel Document Vault, Contextual Moment Engine, AI Travel Concierge, Trip Memory Layer, Travel Insights Dashboard, Corporate Travel Management, Advanced Analytics.

Future vision, explicitly deferred: AI Voice Assistant, Multi-Currency Wallet, AR/VR Destination Preview, Offline AI Assistant, Carbon Impact Tracker.



CROSS-CUTTING: TRUST & TRANSPARENCY LAYER

Not tied to a single release — explainable AI, transparent pricing, live booking confirmation, privacy & consent controls, UPI mandate transparency and an auditable AI decision history run underneath both releases.

AI PRODUCT STRATEGY

Generative AI — conversation, intent, explanation



Agentic AI — research, consensus, budget, booking, support, memory

► Six specialized agents

 Research Agent (Yatra Discover) Finds destinations, flights, hotels, activities → personalized discovery without manual search.	 Consensus Agent (Yatra Sync) Aggregates preferences, resolves conflicts → faster group decisions, fewer disagreements.
 Budget Agent (Yatra Wallet) Estimates and tracks cost, manages the Trip Pool → complete financial transparency.	 Booking Agent (Yatra Book) Monitors pricing, executes bookings on approval → optimized cost, reduced manual effort.
 Support Agent (Yatra Assist) Detects disruptions, rebooks, notifies → reduced stress, higher retention.	 Memory Agent (Yatra Memories) Learns preferences across trips → increasingly personalized recommendations.

MULTI-AGENT WORKFLOW

Specialist agents run in parallel and pass validated context forward; the human approves before anything financial happens.

- 1 User request**
Group shares intent, budget and preferences — in natural language, inside WhatsApp.
- 2 Orchestration**
A shared orchestration layer interprets intent and delegates tasks to specialist agents.
- 3 Parallel specialist agents**
Research, Consensus and Budget agents work in parallel, each contributing context the others can use.
- 4 Validation / context**
Outputs are checked for consistency and confidence before being surfaced to the group.
- 5 Itinerary synthesis**
A consolidated, explainable recommendation is assembled.
- 6 Human approval**
The group reviews cost and plan — a hard gate, not a formality.
- 7 Booking / action**
Booking Agent executes only after approval; funds move only once a booking is verified.
- 8 Post-booking support**
Support Agent monitors the trip and handles disruptions.
- 9 Memory / feedback**
Memory Agent stores the outcome, improving the next trip's recommendations.

**HUMAN APPROVAL IS THE HARD CONSTRAINT**

Every workflow in the source architecture treats booking and payment as actions that require explicit user approval or a pre-defined authorization rule — the AI recommends; the human decides.

AI ARCHITECTURE & TRUST GUARDRAILS



► Responsible AI & governance

- **Human oversight:** bookings, payments and itinerary confirmations always require user approval or pre-defined authorization.
- **Explainable AI:** every recommendation comes with a clear reason — why this destination, this hotel, this itinerary.
- **Privacy & consent:** preferences, history, documents and payment data are encrypted and processed only with explicit consent.
- **Bias monitoring:** recommendation models are evaluated to keep suggestions fair, not commercially skewed.
- **Fallback & recovery:** low AI confidence or an unavailable external system triggers a graceful alternative, never an unreliable answer.
- **Auditability:** an AI decision history gives every recommendation and action a traceable record.

✈️ WHY THIS MATTERS SPECIFICALLY FOR YATRAAI

The platform handles money (UPI Trip Pool), bookings, personal documents (passports, visas), live location, and group-level preference data — a materially higher trust bar than a single-user AI assistant.

VALIDATION STRATEGY

Stage	Question	Key metric / method
1 · Problem	Are we solving the right problem?	JTBD interviews, persona & assumption validation
2 · UX	Can users use the product?	Task completion rate, SUS score, satisfaction
3 · Functional	Does every feature work?	Pass/fail scenario coverage across core workflows
4 · AI & recommendation	Can users trust AI output?	Relevance, accuracy, hallucination rate, confidence communication
5 · Multi-agent	Do agents collaborate correctly?	Orchestration success, context preservation, error recovery
6 · External integration	Do external systems respond reliably?	API latency, retry & fallback behavior
7 · Security & privacy	Is user data protected?	Encryption, consent flows, UPI mandate authorization
8 · Beta	Does it solve real problems?	NPS, CSAT, AI recommendation acceptance
9 · Performance	Can it scale?	Response time, AI inference latency, availability
10 · Continuous monitoring	How do we keep improving?	Consensus completion rate, retention, AI error rate
 THE STAGE WITH NO PLAYBOOK TO BORROW Stage 4 — AI & Recommendation Validation — has no traditional-software equivalent. It is the stage this team can least borrow an existing testing playbook for, and the one worth the most attention.		

PRODUCT DEVELOPMENT & ROADMAP

~26 weeks INDICATIVE TIMELINE	5 BUILD PHASES	28 TASKS SCHEDULED	6–10 ASSUMED TEAM SIZE
---	--------------------------	------------------------------	----------------------------------

P0	Discovery & shaping · wks 1–4 Shape the Consensus Engine algorithm, lock orchestration framework, begin RBI legal consultation in parallel.
P1	Foundation build · wks 5–10 User/group management, WhatsApp Business API integration, core data model, AI preference collection.
P2	Core intelligence · wks 11–18 Research + Consensus + Itinerary Generator, Budget Dashboard, mock booking flow.
P3	Validation & beta · wks 19–22 Functional, AI-quality, multi-agent and security validation; closed beta with friends/family.
P4	Launch readiness · wks 23–26 Performance testing, production monitoring, go/no-go against RBI Payment Aggregator licensing status.




THE DEPENDENCY THAT PROTECTS THE TIMELINE


The RBI Payment Aggregator / escrow authorization is tracked as a parallel legal workstream from Phase 0 onward — never a gate on engineering. That single sequencing choice is what keeps a 6-month MVP plan realistic despite a compliance process the team doesn't control.

PROTOTYPE & UX


Low-fidelity wireframes, screens 1–10 of the end-to-end journey — shown as designed, not described.



01–03 · WhatsApp entry → trip setup → group invite



05–07 · Anonymous consensus vote → destination catalogue → review & propose



08–10 · Full itinerary → UPI pool (held, not spent) → human approval gate before booking

► PM rationale behind the key UX choices

- **WhatsApp-first:** zero install friction, meets travelers where they already coordinate trips.
- **Progressive disclosure:** one question at a time, not a long form.
- **Collaborative decision-making:** anonymous voting removes social pressure from group choices.
- **Trust & transparency:** “held, not spent” wallet state is shown before funds move.
- **Human approval gates:** booking requires an explicit confirm — visible, not buried in settings.

An interactive high-fidelity prototype and a narrated product demonstration video complement these wireframes: euphonious-sherbet-2b8637.netlify.app (prototype link as provided in the source document).

EXPERIMENTATION & WHAT WE LEARNED

Experiment	Method & participants	Key signal	Product implication
Task 01 — JTBD Validation	5 video-recorded 1:1 interview	Fragmented, effortful planning was consistently acknowledged as a real problem	Confirmed the problem statement before any feature was scoped
Task 02 — Two-Channel Onboarding	5 participants, within-subject: WhatsApp vs in-app PWA	WhatsApp: 100% completion, faster, zero install friction. PWA: stronger trust at high-stakes moments (payment, scale)	Hybrid launch: WhatsApp-native onboarding; deep-link to PWA at payment confirmation and for groups of 6+
Task 03 — Renewal Objection Research	5 interviews on 3-month renewal intent	Most common driver: functional & trust issues in AI-driven decisions over time	Priorities Consensus Engine explainability and transparent pricing as retention levers, not just acquisition features
Task 04 — Customer Support Research	3 business-owner interviews (fashion D2C, vehicle rental, home catering)	Fragmented support systems and lack of proactive communication were the top cross-industry pain points	Shaped Yatra Companion™ as proactive and context-aware, not reactive ticket-based support



WHAT TIES THE FOUR EXPERIMENTS TOGETHER

Every experiment tested a specific product assumption — not a general concept — and each produced a concrete design or sequencing decision that shows up elsewhere in this case study.

BUSINESS MODEL & GO-TO-MARKET



- **Primary entry segment:** Urban Group Travelers (22–35, Tier-1 metros) — highest concentration of the consensus pain point, shortest adoption cycle, strongest built-in virality (every group invite is a distribution event).
- **Acquisition approach:** WhatsApp-first, with creator-led distribution (travel influencers, retreat hosts) and content SEO around “group trip planning India.”
- **Key messaging:** “Plan Together. Travel Smarter.” — the only travel product solving group consensus, not just individual booking.
- **Monetization logic:** OTA-style commission (2–4% per booking, primary Year-1 stream), ₹299/month premium subscription, B2B corporate seat licensing, and a 0.5% UPI Pool transaction fee above ₹50K.

⚠ NUMBERS TO LABEL AS ILLUSTRATIVE

India's online travel market ($\approx$ USD 25.38B in 2026, 8.74% CAGR) and MakeMyTrip/ixigo share figures are external market-sizing data, not YatraAI's own validated results. Target metrics such as a 65%+ Group Consensus Rate and a K-factor ≥ 1.2 are product hypotheses to be tested, not achieved outcomes.

RETENTION, SERVICE & ETHICS



► Retention mechanisms

- **Trip Memory as re-engagement:** post-trip recap prompts the next trip, instead of a generic push notification.
- **Personalisation lock-in:** the Memory Agent gets measurably better each trip — a group's third trip is genuinely easier to plan than their first.
- **Visible loyalty signals:** consensus score and trust score give groups a reason to keep improving a number they've started to care about.

► Four ethical risk areas

 Privacy vs. personalisation One person's consent screen shouldn't decide data flow for the whole group. Response: a tiered, Balanced Opt-In Model.	 Bias The Consensus Agent's trade-off logic must not systematically favor louder or higher-spending group members.
 Transparency Every AI recommendation needs a visible reason — explainability is a trust requirement, not a nice-to-have.	 Accountability An auditable AI decision history so a disputed booking or recommendation can always be traced.

✈️ WHY THIS MATTERS HERE SPECIFICALLY

YatraAI handles group-level data, money, personal identity documents, and live location — a materially higher ethical bar than a single-traveller booking engine.

WHAT I WOULD BUILD, VALIDATE & MEASURE NEXT**► What we validated**

- Group travel planning is genuinely fragmented and painful — confirmed across secondary research and 3 primary experiments.
- WhatsApp is the right acquisition channel; the PWA is needed once trust and scale matter.
- Trust and AI reliability, not price, are the strongest renewal risks.

► What remains an assumption

- Group Consensus Rate (65%+) and K-factor (≥ 1.2) are targets, not measured outcomes.
- Live booking, vernacular support and full payment automation are all deferred past Release 1.0 — the MVP is intentionally mocked in these areas.
- RBI Payment Aggregator / escrow authorization timing is outside the team's control.

► What I would test next

- Re-run the two-channel onboarding experiment at 10+ participants to confirm the PWA's scale advantage under real load.
- A live A/B test of the Consensus Engine's explanation copy against actual consensus completion rate.
- Willingness-to-pay validation for the ₹299/month premium tier before it's built.
- A security/consent walkthrough of the Balanced Opt-In Model with real group data.

► What I would measure at launch

- Activation · Consensus completion rate · Time-to-itinerary · Recommendation acceptance · Booking conversion · Budget adherence · Support resolution time · Retention

► What I learned as a PM

Leading the product strategy, discovery and GTM workstream on a multi-agent AI product meant constantly translating between what the AI architecture could technically do and what a group of travelers would actually trust it to do unsupervised. The clearest lesson: every autonomous capability needed a human-approval counterpart before it earned a place in the roadmap — prioritization discipline mattered more here than feature ambition.

"YatraAI is not another booking interface. It is an attempt to coordinate fragmented travel decisions through contextual AI, specialized agents, and human-controlled execution."